

# Applied Bayesian Inference in R Using MCMCpack

Andrew D. Martin

Washington University in St. Louis

Kevin Quinn

Harvard University

<http://mcmcpack.wustl.edu>

June 16, 2006

# MCMCpack Goals

- Free, open-source, easy-to-use software for Bayesian inference.
- Provide a development environment for easy implementation of non-standard statistical models.
- Provide a distribution mechanism for other researchers with a consistent user interface and documentation.

# MCMCpack Design

- “R-like” user interface for Bayesian tools.
- Model-specific design.
- Most computation done in C++ (using the Scythe Statistical Library).
- coda `mcmc` objects for posterior sample storage and summarization.
- Modular design and hidden functions to ease implementing additional models.

## Why Not WinBUGS, JAGS, OpenBugs, etc?

- The BUGS language is good for many things, including quickly developing models. We see it as a complementary tool to MCMCpack.

However . . .

However . . .

- . . . as noted in the WinBUGS manual:

[P]otential users are reminded to be extremely careful if using this program for *serious statistical analysis*. . . If there is a problem, WinBUGS might just crash, which is not very good, but it might well carry on and produce answers that are wrong, which is even worse (p. 1).

However . . .

- . . . as noted in the WinBUGS manual:

[P]otential users are reminded to be extremely careful if using this program for *serious statistical analysis*. . . If there is a problem, WinBUGS might just crash, which is not very good, but it might well carry on and produce answers that are wrong, which is even worse (p. 1).

- The BUGS language can be slow (especially for large problems), and the WinBUGS engine does not work for certain problems.

However . . .

- . . . as noted in the WinBUGS manual:

[P]otential users are reminded to be extremely careful if using this program for *serious statistical analysis*. . . If there is a problem, WinBUGS might just crash, which is not very good, but it might well carry on and produce answers that are wrong, which is even worse (p. 1).

- The BUGS language can be slow (especially for large problems), and the WinBUGS engine does not work for certain problems.
- The various flavors of BUGS require the user to do some modest programming which increases the opportunities for something to go wrong.



# Example 1: Inference for an Item Response Theory (IRT) Model

Consider:

$$y_{ij} \sim \text{Bernoulli}(\pi_{ij}), \quad i = 1, \dots, I \quad j = 1, \dots, J$$

with

$$\pi_{ij} = \Phi(-\alpha_j + \beta_j \theta_i)$$

Typically,  $i$  indexes test-takers (subjects) and  $j$  indexes test items.

Prior distributions for the model parameters are:

$$(\alpha_j, \beta_j)' \sim \mathcal{N}(\mathbf{a}_0, \mathbf{A}_0^{-1}) \quad j = 1, \dots, J$$

and

$$\theta_i \sim \mathcal{N}(0, 1) \quad i = 1, \dots, I$$

To fit this model to some test data from  
<http://work.psych.uiuc.edu/irt/downloads.asp>  
using MCMCpack one can do the following.

```
post.irt <- MCMCirt1d(testdata,  
                      burnin=5000, mcmc=100000, thin=10,  
                      AB0=.001, store.item=TRUE, store.ability=TRUE,  
                      verbose=5000,  
                      theta.constraints=list(SUBJECT368="+",  
                                             SUBJECT356="-"))
```









To fit this model to some test data from <http://work.psych.uiuc.edu/irt/downloads.asp> using MCMCpack one can do the following.

```
post.irt <- MCMCirt1d(testdata,  
                      burnin=5000, mcmc=100000, thin=10,  
                      AB0=.001, store.item=TRUE, store.ability=TRUE,  
                      verbose=5000,  
                      theta.constraints=list(SUBJECT368="+",  
                                             SUBJECT356="-"))
```

To look at some traceplots and marginal density estimates one can use:

```
plot(post.irt)
```

which produces:



Subject 436 and subject 447 both got 75% of the test items correct. Which subject has higher ability?

The posterior probability that subject 436 has higher ability than subject 447 can be calculated with:

```
mean(post.irt[,436] > post.irt[,447])
```

which evaluates to 0.65.

We might also be interested in which test items do a good job of discriminating between high and low ability subjects. The  $\beta$  parameters provide this information.

We can calculate the posterior expectation of each of the discrimination parameters with:

```
beta.post <- post.irt[,seq(from=452, to=530, by=2)]  
print(sort(colMeans(beta.post)))
```

which produces:

beta.ITEM18	beta.ITEM28	beta.ITEM16	beta.ITEM37	beta.ITEM19	beta.ITEM32
-0.033332789	0.007374421	0.084904972	0.217837674	0.276428999	0.315174596
beta.ITEM8	beta.ITEM20	beta.ITEM40	beta.ITEM21	beta.ITEM1	beta.ITEM34
0.331350234	0.411436261	0.418098884	0.419030043	0.439532906	0.460505549
beta.ITEM7	beta.ITEM10	beta.ITEM22	beta.ITEM14	beta.ITEM35	beta.ITEM31
0.462356437	0.472680964	0.481691248	0.482167728	0.514685111	0.564987099
beta.ITEM39	beta.ITEM36	beta.ITEM33	beta.ITEM3	beta.ITEM4	beta.ITEM27
0.601600337	0.623693132	0.627042964	0.634003442	0.641052025	0.672391363
beta.ITEM29	beta.ITEM2	beta.ITEM6	beta.ITEM11	beta.ITEM30	beta.ITEM13
0.689009747	0.706951188	0.723142300	0.762149494	0.772004859	0.779142756
beta.ITEM9	beta.ITEM38	beta.ITEM5	beta.ITEM23	beta.ITEM25	beta.ITEM17
0.851232006	0.901117855	0.902609937	0.969587117	1.039535127	1.041296658
beta.ITEM12	beta.ITEM26	beta.ITEM24	beta.ITEM15		
1.099500195	1.217643770	1.282921879	1.377932523		



We can also calculate the posterior expected ranks of the discrimination parameters:

```
beta.post.ranks <- matrix(NA, nrow(beta.post), ncol(beta.post))
colnames(beta.post.ranks) <- colnames(beta.post)
for (i in 1:nrow(beta.post)){
  beta.post.ranks[i,] <- rank(beta.post[i,])
}
print(sort(colMeans(beta.post.ranks)))
```

beta.ITEM18	beta.ITEM28	beta.ITEM16	beta.ITEM37	beta.ITEM19	beta.ITEM32
1.4852	1.9182	2.8395	4.7514	5.8596	7.0272
beta.ITEM8	beta.ITEM20	beta.ITEM21	beta.ITEM40	beta.ITEM1	beta.ITEM34
7.5302	10.7339	11.1222	11.1408	12.1650	13.1280
beta.ITEM7	beta.ITEM10	beta.ITEM22	beta.ITEM14	beta.ITEM35	beta.ITEM31
13.2174	13.6764	14.0879	14.1981	15.8325	18.3422
beta.ITEM39	beta.ITEM36	beta.ITEM33	beta.ITEM3	beta.ITEM4	beta.ITEM27
20.1938	21.2814	21.4698	21.7877	22.1645	23.6409
beta.ITEM29	beta.ITEM2	beta.ITEM6	beta.ITEM11	beta.ITEM30	beta.ITEM13
24.4503	25.2602	25.9368	27.5372	27.9766	28.2097
beta.ITEM9	beta.ITEM38	beta.ITEM5	beta.ITEM23	beta.ITEM25	beta.ITEM17
30.6817	32.0941	32.1787	33.7927	35.0096	35.2362
beta.ITEM12	beta.ITEM26	beta.ITEM24	beta.ITEM15		
36.2589	37.8847	38.5471	39.3517		

## Example 2: Calculating Bayes Factors for Model Comparison

Suppose that the observed data  $\mathbf{y}$  could have been generated under one of two models  $\mathcal{M}_1$  and  $\mathcal{M}_2$ .

A natural thing to ask from the Bayesian perspective is, “what is the posterior probability that  $\mathcal{M}_1$  is true (assuming either  $\mathcal{M}_1$  or  $\mathcal{M}_2$  is true)?”

How can we calculate this? Well, using Bayes theorem we can write:

$$\Pr(\mathcal{M}_k|\mathbf{y}) = \frac{p(\mathbf{y}|\mathcal{M}_k) \Pr(\mathcal{M}_k)}{p(\mathbf{y}|\mathcal{M}_1) \Pr(\mathcal{M}_1) + p(\mathbf{y}|\mathcal{M}_2) \Pr(\mathcal{M}_2)}, \quad k = 1, 2$$

The key quantities here are  $p(\mathbf{y}|\mathcal{M}_1)$  and  $p(\mathbf{y}|\mathcal{M}_2)$  which are called *marginal likelihoods* or *integrated likelihoods*.

It is instructive to look at the posterior odds in favor of one model (say  $\mathcal{M}_1$ ):

$$\frac{\Pr(\mathcal{M}_1|\mathbf{y})}{\Pr(\mathcal{M}_2|\mathbf{y})} = \frac{p(\mathbf{y}|\mathcal{M}_1)}{p(\mathbf{y}|\mathcal{M}_2)} \times \frac{\Pr(\mathcal{M}_1)}{\Pr(\mathcal{M}_2)}$$

What this means is that if we want to move from the prior odds in favor of  $\mathcal{M}_1$  to the posterior odds in favor of  $\mathcal{M}_1$  we simply multiply the prior odds by:

$$B_{12} = \frac{p(\mathbf{y}|\mathcal{M}_1)}{p(\mathbf{y}|\mathcal{M}_2)}$$

which is called the *Bayes factor* for  $\mathcal{M}_1$  relative to  $\mathcal{M}_2$ .

In words,

$$\text{posterior odds} = \text{Bayes factor} \times \text{prior odds}$$

Since the posterior odds equal the Bayes factor when the models are equally likely *a priori*, the Bayes factor is a measure of how much support is available in the data for one model relative to another.

It is now possible to calculate Bayes factors and posterior probabilities of models with MCMCpack.

```
data(birthwt)
model1 <- MCMCregress(bwt~age+lwt+as.factor(race) + smoke + ht,
                      data=birthwt,
                      b0=c(2700, 0, 0, -500, -500, -500, -500),
                      B0=c(1e-6, .01, .01, 1.6e-5, 1.6e-5, 1.6e-5,
                          1.6e-5), c0=10, d0=4500000,
                      marginal.likelihood="Chib95", mcmc=50000)

model2 <- MCMCregress(bwt~age+lwt+as.factor(race) + smoke,
                      data=birthwt,
                      b0=c(2700, 0, 0, -500, -500, -500),
                      B0=c(1e-6, .01, .01, 1.6e-5, 1.6e-5, 1.6e-5),
                      c0=10, d0=4500000,
                      marginal.likelihood="Chib95", mcmc=50000)
```

```
model3 <- MCMCregress(bwt~as.factor(race) + smoke + ht,  
                     data=birthwt,  
                     b0=c(2700, -500, -500, -500, -500),  
                     B0=c(1e-6, 1.6e-5, 1.6e-5, 1.6e-5,  
                          1.6e-5), c0=10, d0=4500000,  
                     marginal.likelihood="Chib95", mcmc=50000)  
  
BF <- BayesFactor(model1, model2, model3)  
  
mod.probs <- PostProbMod(BF)
```

```
print(BF)
```

The matrix of Bayes Factors is:

	model1	model2	model3
model1	1.000	14.08	7.289
model2	0.071	1.00	0.518
model3	0.137	1.93	1.000

The matrix of the natural log Bayes Factors is:

	model1	model2	model3
model1	0.00	2.645	1.986
model2	-2.64	0.000	-0.658
model3	-1.99	0.658	0.000

```
print(mod.probs)
```

model1	model2	model3
0.82766865	0.05878317	0.11354819



## Example 3: Single Block Metropolis Sampling for a User-Defined Model

MCMCpack also has some facilities for fitting user-specified models.

The `MCMCmetrop1R()` function uses a random walk Metropolis algorithm to sample from a user-defined log-posterior density.

Suppose one is interested in fitting a Bayesian negative binomial regression to the Ornstein data in the `car` package.

One could do the following:

```

negbinlogpost<- function(theta, y, X, b0, B0, c0, d0){

  ## log of inverse gamma density
  logdinvgamma <- function(sigma2, a, b){
    logf <- a * log(b) - lgamma(a) + -(a+1) * log(sigma2) + -b/sigma2
    return(logf)
  }

  ## log of multivariate normal density
  logdmvnorm <- function(theta, mu, Sigma){
    d <- length(theta)
    logf <- -0.5*d * log(2*pi) - 0.5*log(det(Sigma)) -
      0.5 * t(theta - mu) %*% solve(Sigma) %*% (theta - mu)
    return(logf)
  }

  k <- length(theta)
  beta <- theta[1:(k-1)]
  alpha <- exp(theta[k])
  mu <- exp(X %*% beta)

  ## evaluate log-likelihood at (alpha, beta)
  log.like <- sum(
    lgamma(y+alpha) - lfactorial(y) - lgamma(alpha) +

```

```

        alpha * log(alpha/(alpha+mu)) +
        y * log(mu/(alpha+mu))
    )

## evaluate log prior at (alpha, beta)
## note Jacobian term necessary b/c of transformation
log.prior <- logdinvgamma(alpha, c0, d0) + theta[k] +
  logdmvnorm(beta, b0, B0)

  return(log.like+log.prior)
}

library(car)
data(Ornstein)
yvec <- Ornstein$interlocks
Xmat <- model.matrix(~sector+nation, data=Ornstein)

post.negbin <- MCMCmetrop1R(negbinlogpost, theta.init=rep(0,14),
                           X=Xmat, y=yvec,
                           thin=2, mcmc=50000, burnin=1000,
                           tune=rep(.7,14),
                           verbose=500, logfun=TRUE, optim.maxit=100,
                           b0=0, B0=diag(13)*1000, c0=1, d0=1)

```

```
summary(post.negbin)
```

```
Iterations = 1001:50999
```

```
Thinning interval = 2
```

```
Number of chains = 1
```

```
Sample size per chain = 25000
```

1. Empirical mean and standard deviation for each variable,  
plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
beta.(Intercept)	2.34860	0.1602	0.0010132	0.003964
beta.sectorBNK	1.70677	0.4014	0.0025385	0.010942
beta.sectorCON	-0.59210	0.5526	0.0034949	0.015653
beta.sectorFIN	0.99221	0.2653	0.0016778	0.007198
beta.sectorHLD	0.30736	0.4251	0.0026884	0.011383
beta.sectorMAN	0.12265	0.2226	0.0014077	0.005949
beta.sectorMER	0.23269	0.2728	0.0017253	0.007095
beta.sectorMIN	0.66688	0.2121	0.0013413	0.006118
beta.sectorTRN	0.89300	0.2779	0.0017578	0.006900
beta.sectorWOD	0.67976	0.2752	0.0017403	0.008080
beta.nationOTH	-0.07207	0.2985	0.0018882	0.008870
beta.nationUK	-0.56004	0.2736	0.0017304	0.008131
beta.nationUS	-0.84186	0.1511	0.0009559	0.004240

```
log(alpha)          0.07750 0.1070 0.0006764          0.003525
```

2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
beta.(Intercept)	2.0484	2.23548	2.34587	2.4571	2.668864
beta.sectorBNK	0.9810	1.42833	1.68293	1.9583	2.555186
beta.sectorCON	-1.6220	-0.96827	-0.59939	-0.2372	0.538723
beta.sectorFIN	0.4877	0.81116	0.98990	1.1638	1.539247
beta.sectorHLD	-0.4815	0.02123	0.29320	0.5806	1.182072
beta.sectorMAN	-0.3233	-0.02640	0.12499	0.2735	0.557485
beta.sectorMER	-0.2974	0.05259	0.22863	0.4068	0.776510
beta.sectorMIN	0.2526	0.52381	0.66726	0.8065	1.088377
beta.sectorTRN	0.3814	0.70355	0.88244	1.0776	1.448740
beta.sectorWOD	0.1624	0.49518	0.67317	0.8621	1.244756
beta.nationOTH	-0.6499	-0.27512	-0.07392	0.1272	0.512434
beta.nationUK	-1.0953	-0.74036	-0.56804	-0.3831	-0.005212
beta.nationUS	-1.1399	-0.94211	-0.84336	-0.7415	-0.539916
log(alpha)	-0.1357	0.00676	0.07577	0.1511	0.287035

## Future Plans

- add more models (generalized linear mixed models, Dirichlet process mixture models, dynamic linear models, other models suggested by users)
- add more options for prior specifications
- add additional flexible sampling engines (more complicated Metropolis-Hastings methods, slice sampling, etc.)
- improve performance
- add vignettes
- add enhancements to coda mcmc objects to allow for sampling from posterior and prior predictive distributions

# Acknowledgements

National Science Foundation  
Program in Methodology, Measurement, and Statistics  
(Grants SES-0350646 and SES-0350613)

Washington University, Department of Political Science, and the  
Weidenbaum Center on the Economy, Government, and Public Policy  
<http://wc.wustl.edu>

Harvard University, Department of Government, and the  
Institute for Quantitative Social Science  
<http://www.iq.harvard.edu>